

Неструктурирани бази на податоци  
Компаративна анализа на  
**Релациони и Граф-базиранани**  
бази на податоци

Марија Вецовска

185008

## Вовед

Целта на овој проект е да се направи компаративна анализа на две бази на податоци од различен тип – неструктурирана и релациона. Анализата се состои од извршување на прашалници врз истите податоци во двете бази и споредба на брзината на извршување.

Податочното множество е Identifying Influential Bloggers: Techcrunch и е преземено од Kaggle.

Множеството е достапно на линкот: <https://www.kaggle.com/datasets/lakritidis/identifying-influential-bloggers-techcrunch>.

Неструктурирана база на податоци која се користи во рамките на овој проект е Граф-базираната база на податоци Neo4j, а релациона база е базата PostgreSQL.

### Податочното множество:

Множеството Identifying Influential Bloggers: Techcrunch се состои од 3 фајлови:

- authors.csv – податоци за 107 блогери на Techcrunch,
- posts.csv – податоци за 19464 блог постови на Techcrunch напишани од авторите,
- comments.csv – коментари од 746561 читатели на постовите
- inlinks.csv – податоци за 193808 страници кои ги линкуваат Techcrunch постовите.

### Чистење на податоците

Во рамки на податоците се среќаваат многу нецелосни податоци, исти редови со различно ИД и други грешки. Од тие причини csv фајловете се трансформираат со помош на python скрипта:

- authors\_transformed.csv: ги содржи сите автори на Блогови, Постови и Коментари (автори од authors.csv, posts.csv и comments.csv).
- posts\_transformed.csv: ги содржи податоците за Блогови и страници кои ги линкуваат (постови од authors.csv и inlinks.csv).
- comments\_transformed.csv: ги содржи податоците за коментари од корисници (коментари од comments.csv).
- inlinks\_transformed.csv: ги содржи линковите меѓу сите постови (линкови од inlinks.csv)

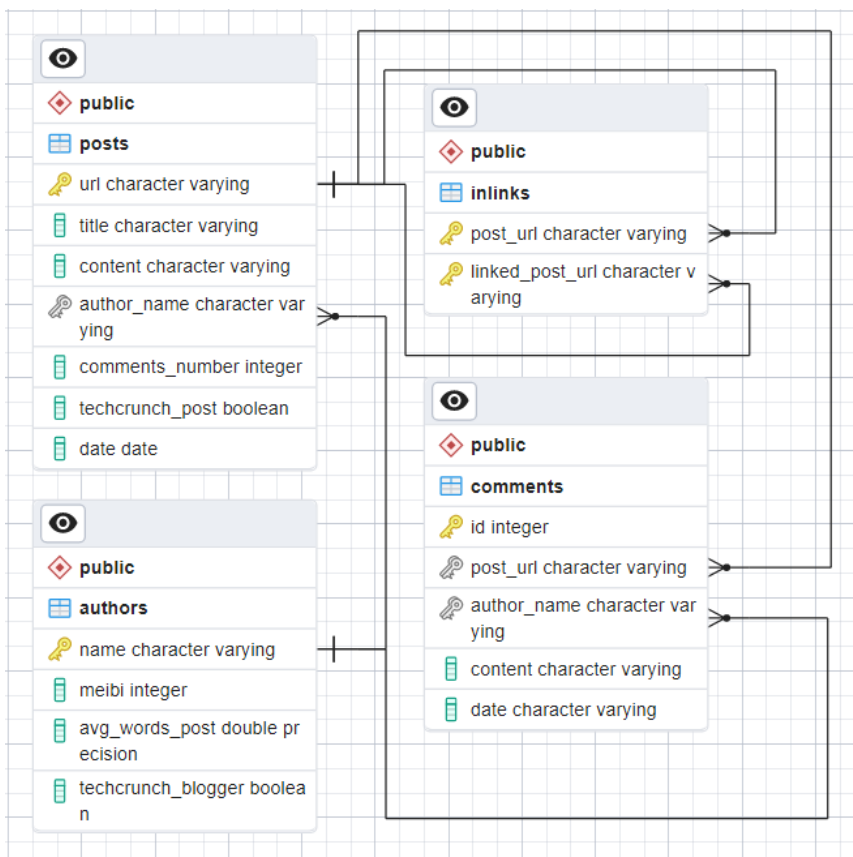
На крај се добиваат следните csv:

```
Authors:
[{'name', 'meibi', 'avg_words_post', 'techcrunch_blogger'}]
[{'Jason Kincaid', '43', '4.257910', 'true'}]
Authors count:196273
Posts:
[{'url', 'title', 'content', 'author_name', 'num_comments', 'date', 'techcrunch_post'}]
[{'http://techcrunch.com/2010/04/01/we-just-tested-twitters-anywhere-platform-screenshots/', 'We Just Tested Twitter6#8217;s @anywhere Platform (Screenshots)', 'During his keynote'}]
Posts count:178305
Comments:
[{'id', 'post_url', 'content', 'author_name', 'date'}]
[{'1', 'http://techcrunch.com/2010/04/01/we-just-tested-twitters-anywhere-platform-screenshots/', 'Seemed to work fine.', 'BJ Cook', '2010-04-01'}]
Comments count:741276
Inlinks:
[{'post', 'linked_post'}]
[{'http://feeds.notaniche.com/~c/Afrison/~3/hXARXfINteM/', 'http://techcrunch.com/2005/06/11/technorati-new-improved/'}]
Inlinks count:193504
```

## Внесување на податоците во PostgreSQL

- Креирање на табели и Шемата на базата на податоци:

```
create table Authors (  
    name varchar primary key,  
    meibi integer,  
    avg_words_post float,  
    techcrunch_blogger boolean not null  
);  
  
create table Posts (  
    url varchar primary key,  
    title varchar,  
    content varchar,  
    author_name varchar,  
    comments_number integer,  
    techcrunch_post boolean not null,  
    date date,  
    foreign key(author_name) references Authors(name)  
);  
  
create table Comments (  
    id integer,  
    post_url varchar,  
    author_name varchar,  
    content varchar,  
    date varchar,  
    foreign key(post_url) references Posts(url),  
    foreign key(author_name) references Authors(name),  
    primary key(id)  
);  
  
create table Inlinks (  
    post_url varchar,  
    linked_post_url varchar,  
    primary key(post_url, linked_post_url),  
    foreign key(post_url) references Posts(url),  
    foreign key(linked_post_url) references Posts(url)  
);
```



## - Внесување на податоците

### ○ Автори

Query	Query History
<pre>1 COPY authors(name, meibi, avg_words_post, techcrunch_blogger) 2 FROM 'C:/Users/marija/Documents/Neo4j/relate-data/dbmss/dbms-535955cc-1561-4ae6-ac9d-cb27f2726a0e/import/authors_transformed.csv' 3 DELIMITER ',' 4 CSV HEADER;</pre>	
Data Output	Messages
COPY 196273	
Query returned successfully in 2 secs 11 msec.	

### ○ Постови

Query	Query History
<pre>1 COPY posts(url, title, content, author_name, comments_number, date, techcrunch_post ) 2 FROM 'C:/Users/marija/Documents/Neo4j/relate-data/dbmss/dbms-535955cc-1561-4ae6-ac9d-cb27f2726a0e/import/posts_transformed.csv' 3 DELIMITER ',' 4 CSV HEADER;</pre>	
Data Output	Messages
COPY 170305	
Query returned successfully in 8 secs 483 msec.	

### ○ Коментари

Query	Query History
<pre>1 COPY comments(id, post_url, content, author_name, date) 2 FROM 'C:/Users/marija/Documents/Neo4j/relate-data/dbmss/dbms-535955cc-1561-4ae6-ac9d-cb27f2726a0e/import/comments_transformed.csv' 3 DELIMITER ',' 4 CSV HEADER;</pre>	
Data Output	Messages
COPY 741276	
Query returned successfully in 48 secs 96 msec.	

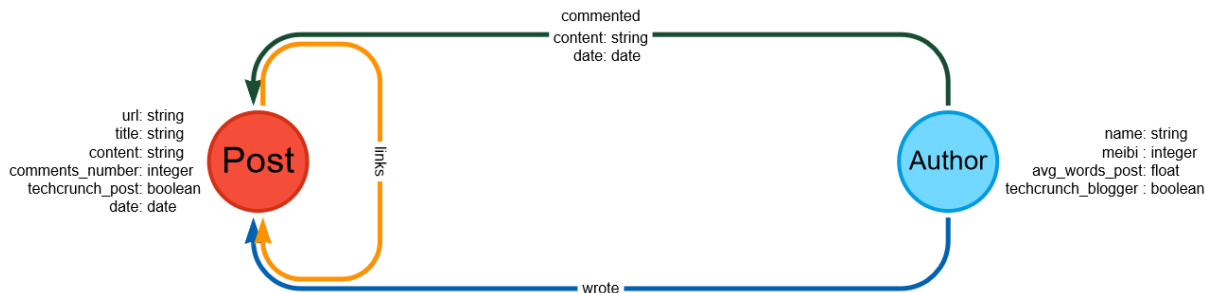
### ○ Линкови

Query	Query History
<pre>1 COPY inlinks(post_url, linked_post_url) 2 FROM 'C:/Users/marija/Documents/Neo4j/relate-data/dbmss/dbms-535955cc-1561-4ae6-ac9d-cb27f2726a0e/import/inlinks_transformed.csv' 3 DELIMITER ',' 4 CSV HEADER;</pre>	
Data Output	Messages
COPY 193504	
Query returned successfully in 16 secs 460 msec.	

**Вкупно време на внесување во база: 75,05 секунди**

## Внесување на податоците во Neo4j

- Шема на базата:



- Подесување на ресурси достапни за базата

### Edit settings

```
# Disable logging JMX endpoint.
server.jvm.additional=-Dlog4j2.disable.jmx=true

# Limit JVM metaspace and code cache to allow garbage collection. Used by cypher for code generation and
# may grow indefinitely unless constrained.
# Useful for memory constrained environments
server.jvm.additional=-XX:MaxMetaspaceSize=1024m
server.jvm.additional=-XX:ReservedCodeCacheSize=512m

# *****
# Wrapper Windows NT/2000/XP Service Properties
# *****
# WARNING - Do not modify any of these properties when an application
# using this configuration file has been installed as a service.
# Please uninstall the service before modifying this section. The
# service can then be reinstalled.

# Name of the service
server.windows_service_name=neo4j

# *****
# Other Neo4j system properties
# *****

dbms.memory.heap.initial_size=3G
dbms.memory.heap.max_size=3G
dbms.memory.pagecache.size=3G
dbms.windows_service_name=neo4j-relate-dbms-535955cc-1561-4ae6-ac9d-cb27f2726a0e
dbms.jvm.additional=-Dlog4j2.formatMsgNoLookups=true
```

Reset to defaults

Undo

Apply

Close

## - Ограничувања

neo4j\$ SHOW CONSTRAINTS

	id	name	type	entityType	labelsOrTypes	properties	ownedIndex
1	4	"author_name"	"UNIQUENESS"	"NODE"	["Author"]	["name"]	"author_name"
2	6	"post_url"	"UNIQUENESS"	"NODE"	["Post"]	["url"]	"post_url"

Started streaming 2 records after 40 ms and completed after 68 ms.

## - Автори

```
1 :auto LOAD CSV WITH HEADERS FROM "file:///authors_transformed.csv" AS row
2 CALL {
3   WITH row
4   MERGE(a:Author{name:row.name})
5   SET a.meibi=toInteger(row.meibi)
6   SET a.avg_words_post=toFloat(row.avg_words_post)
7   SET a.techcrunch_blogger=toBoolean(row.techcrunch_blogger)
8 } IN TRANSACTIONS
```

Added 196273 labels, created 196273 nodes, set 392760 properties, completed after 16115 ms.

Added 196273 labels, created 196273 nodes, set 392760 properties, completed after 16115 ms.

## - Постови

```
1 :auto LOAD CSV WITH HEADERS FROM "file:///posts_transformed.csv" AS row
2 CALL {
3   WITH row
4   MERGE(p:Post{url:row.url})
5   SET p.title=row.title
6   SET p.content=row.content
7   SET p.comments_number=toInteger(row.comments_number)
8   SET p.techcrunch_post=toBoolean(row.techcrunch_post)
9   SET p.date=date(row.date)
10 } IN TRANSACTIONS
```

Added 170305 labels, created 170305 nodes, set 700072 properties, completed after 15970 ms.

Added 170305 labels, created 170305 nodes, set 700072 properties, completed after 15970 ms.

- Релацијата **:wrote** која означува кој е авторот на кој пост

```
1 :auto LOAD CSV WITH HEADERS FROM "file:///posts_transformed.csv" AS row
2 CALL {
3   WITH row
4   MATCH(p:Post{url:row.url})
5   MATCH(a:Author{name:row.author_name})
6   MERGE(a)-[:wrote]-(p)
7 } IN TRANSACTIONS
```

Table

Code

Created 153901 relationships, completed after 10502 ms.

Created 153901 relationships, completed after 10502 ms.

- Релацијата **:commented** која претставува коментар од еден автор на некој пост

```
1 :auto LOAD CSV WITH HEADERS FROM "file:///comments_transformed.csv" AS row
2 CALL {
3   WITH row
4   MATCH(p:Post{url:row.post_url})
5   MATCH(a:Author{name:row.author_name})
6   MERGE(a)-[c:commented{id:row.id, date:row.date}]->(p) SET c.content = row.comment
7 } IN TRANSACTIONS
```

Table

Code

Set 1482552 properties, created 741276 relationships, completed after 63114 ms.

Set 1482552 properties, created 741276 relationships, completed after 63114 ms.

- Релацијата **:links** која означува линк од еден пост према друг

```
1 :auto LOAD CSV WITH HEADERS FROM "file:///inlinks_transformed.csv" AS row
2 CALL {
3   WITH row
4   MATCH(p:Post{url:row.post})
5   MATCH(l:Post{url:row.linked_post})
6   MERGE(p)-[:links]-(l)
7 } IN TRANSACTIONS
```

Table

Code

Created 193504 relationships, completed after 19657 ms.

Created 193504 relationships, completed after 19657 ms.

**Вкупно време на внесување во база: 109.234 минути**

- **Додавање на дополнителни лабели:**

```
1 MATCH (a:Author{techcrunch_blogger:true})
2 SET a:Blogger
3 RETURN a
```



a

1

```
{
  "identity": 46964,
  "labels": [
    "Author",
    "Blogger"
  ],
  "properties": {
    "avg_words_post": 4.25791,
    "meibi": 43,
    "name": "Jason Kincaid",
    "techcrunch_blogger": true
  },
  "elementId": "46964"
}
```

2

```
{
```

Added 107 labels, started streaming 107 records after 3 ms and completed after 391 ms.

```
1 MATCH (p:Post{techcrunch_post:true})
2 SET p:Blog
3 RETURN p
```



p

1

```
{
  "identity": 243237,
  "labels": [
    "Post",
    "Blog"
  ],
  "properties": {
    "date": "2010-04-01",
    "title": "We Just Tested Twitter's @anywhere Platform (Screenshots)",
    "techcrunch_post": true,
    "url": "http://techcrunch.com/2010/04/01/we-just-tested-twitters-anywhere-platform-screen",
    "content": "During his keynote at SXSW last month, Twitter CEO Evan Williams announced an i
looks like we've just come across the first site to feature @anywhere. Meet Eggboiling
Connect With Twitter; buttons to follow twitter users @jack, @biz, and @ev; a t
each of the follow buttons appropriately changed the status from Follow i
Login/Connect to Twitter (but before image/report link have loaded)Tags/Report link..."
  }
}
```

Added 19458 labels, started streaming 19458 records after 3 ms and completed after 70 ms, displaying first 1000 rows.



## Прашалници:

1. Сите techcrunch блогови кои се напишани помеѓу 2009 и 2010 година и во насловот содржат 'twitter'

### Neo4j:

```
1 MATCH(b:Blog)
2 WHERE b.date ≥ date('2009-01-01') AND b.date ≤ date('2010-12-31')
3 AND (b.title CONTAINS 'twitter' OR b.title CONTAINS 'Twitter')
4 RETURN b.title
```

	b.title
1	"We Just Tested Twitter&#8217;s @anywhere Platform (Screenshots)"
2	"Twitter Tweaks Search, A Few Popular Tweets Now At The Top Of Results"
3	"SharesPost Report Values Twitter At \$750 Million"
4	"Twitter Sets Chirp Free&#8230; Well, Cheaper"
5	"Twitter Launches A New, Dynamic Homepage"
6	"Twitter And The Nine-Month Bounce"
7	

Started streaming 767 records in less than 1 ms and completed after 53 ms.

### PostgreSQL:

```
7
8 SELECT p.title
9 FROM posts as p
10 WHERE p.techcrunch_post = 'true'
11 AND p.date>=date('2009-01-01') AND p.date<=date('2010-12-31')
12 AND ( p.title like '%twitter%' OR p.title like '%Twitter%')
13
```

	title
	character varying
1	We Just Tested Twitter&#8217;s @anywhere Platform (Screenshots)
2	Twitter Tweaks Search, A Few Popular Tweets Now At The Top Of Results
3	SharesPost Report Values Twitter At \$750 Million
4	Twitter Sets Chirp Free&#8230; Well, Cheaper
5	Twitter Launches A New, Dynamic Homepage
6	Twitter And The Nine-Month Bounce

Total rows: 767 of 767    Query complete 00:00:00.107

Разлика: 54 ms.

## 2. Постовите напишани од автор кои содржат линк кон пост од истиот автор

### Neo4j:

```
1 MATCH(p1:Post)-[:links]→(p)
2 MATCH (a:Author)-[:wrote]→(p)
3 MATCH (a)-[:wrote]→(p1)
4 RETURN distinct(p.url)
5
```

	(p.url)
1	"http://www.mobilecrunch.com/2010/02/17/opera-says-opera-mini-for-iphone-100-compliant-with-app-store-policies-video/"
2	"http://techcrunch.com/2009/12/16/google-android-market/"
3	"http://techcrunch.com/2010/03/05/global-smartphone-app-download-market-could-reach-15-billion-by-2013-report/"
4	"http://techcrunch.com/2010/02/28/19-startups-showing-their-wares-at-techcrunch-japans-tokyocamp-demo-event/"
5	"http://techcrunch.com/2010/02/15/layar-scores-3-4m-in-funding-global-distribution-agreement/"
6	"http://techcrunch.com/2009/10/14/layar-brings-augmented-reality-browser-to-the-iphone-screenshots/"
7	

Started streaming 42 records after 1 ms and completed after 446 ms.

### PostgreSQL:

```
8 SELECT p.url
9 FROM posts as p
10 JOIN inlinks as il ON il.post_url=p.url
11 JOIN posts as p1 on il.linked_post_url = p1.url
12 JOIN authors as a ON p.author_name = a.name
13 WHERE p1.author_name=a.name
14
```

Data Output Messages Notifications

	url [PK] character varying
1	http://www.mobilecrunch.com/2008/12/09/android-rising-sony-ericsson-vodafone-and-14-others-join-open-handset-alliance/
2	http://www.mobilecrunch.com/2009/01/28/review-t-mobile-blackberry-8900/
3	http://www.mobilecrunch.com/2009/03/20/sekai-camera-mobile-social-tagging-is-coming-to-android-phones-too/
4	http://www.mobilecrunch.com/2010/03/03/skype-for-symbian-lands-on-ovi-store-more-than-200-million-possible-users/
5	http://www.mobilecrunch.com/2009/05/05/att-finally-releases-what-should-have-been-the-first-iphone-app/
6	http://www.crunchgear.com/2009/04/22/proof-that-att-needs-to-extend-that-ipl

Total rows: 48 of 48 Query complete 00:00:00.266

✓ Successfully run. Total query runtime: 266

Разлика: - 180 ms

- Блоговите напишани од некој автор кои содржат линк кон пост од истиот автор, кој има барем 30 коментари.

#### Neo4j:

```
1 MATCH(p1:Blog)-[:links]→(p)
2 WHERE p:Blog
3 MATCH (a:Blogger)-[:wrote]→(p)
4 MATCH (a)-[:wrote]→(p1)
5 MATCH (p)←[c:commented]-(c)
6 WITH p, COUNT(c) AS num
7 WHERE num >= 30
8 RETURN p.url, num
```

	p.url	num
1	"http://techcrunch.com/2007/11/05/breaking-google-announces-android-and-open-handset-alliance/"	112
2	"http://techcrunch.com/2008/03/17/microsoft-adopts-flash-lite-for-windows-mobile-as-a-stopgap-measure/"	52
3	"http://techcrunch.com/2008/04/30/adobes-open-screen-project-write-once-flash-everywhere/"	66
4	"http://www.mobilecrunch.com/2008/11/16/adobe-to-demo-flash-on-mobile-but-only-windows-still-working-on-the-iphone/"	30
5	"http://www.mobilecrunch.com/2009/07/20/review-htc-hero/"	73
6	"http://www.crunchgear.com/2009/09/17/review-htc-hero-from-sprint/"	92
7		

Started streaming 27 records after 1 ms and completed after 42 ms.

#### PostgreSQL:

```
9 SELECT distinct(p1.url)
10 FROM posts as p
11 JOIN inlinks as il ON il.post_url=p.url AND p.techcrunch_post IS TRUE
12 JOIN posts as p1 on il.linked_post_url = p1.url AND p1.techcrunch_post = 'true'
13 JOIN authors as a ON p.author_name = a.name AND p1.author_name=a.name
14 JOIN comments as c on c.post_url = p1.url
15 GROUP BY p1.url
16 HAVING count(c.id) >= 30
17
```

	url
1	http://techcrunch.com/2007/11/05/breaking-google-announces-android-and-open-handset-alliance/
2	http://techcrunch.com/2008/03/17/microsoft-adopts-flash-lite-for-windows-mobile-as-a-stopgap-measure/
3	http://techcrunch.com/2008/04/30/adobes-open-screen-project-write-once-flash-everywhere/
4	http://techcrunch.com/2009/02/17/tagging-the-real-world-sekai-camera-for-the-iphone-is-alive-and-very-cool/
5	http://techcrunch.com/2009/04/02/skype-iphone-app-downloaded-one-million-times-in-first-two-days/
6	http://techcrunch.com/2009/04/14/shocker-att-wants-to-keep-keep-sitting-on-its-golden-egg-the-iphone/

Total rows: 27 of 27    Query complete 00:00:00.274

Разлика: 232 ms

#### 4. Број на низи од 4 поста поврзани со линкови

##### Neo4j:

```
neo4j$ MATCH(p:Post)-[:links *3]→(p1:Post) RETURN count(p)
```

	count(p)
1	295

Started streaming 1 records after 1 ms and completed after 297 ms.

##### PostgreSQL:

```
9 SELECT count(p)
10 FROM posts as p
11 JOIN inlinks as il ON il.post_url=p.url
12 JOIN posts as p1 on il.linked_post_url = p1.url
13 JOIN inlinks as il1 ON il1.post_url=p1.url
14 JOIN posts as p2 on il1.linked_post_url = p2.url
15 JOIN inlinks as il2 ON il2.post_url=p2.url
16 JOIN posts as p3 on il2.linked_post_url = p3.url
17
```

Data Output Messages Notifications

	count bigint
1	297

Total rows: 1 of 1 Query complete 00:00:00.734

✓ Successfully run. Total time: 00:00:00.734

Разлика : 437 мс

## 5. Број на автори кои имаат напишано и techcrunch и обичен пост

### Neo4j:

The image shows the Neo4j Cypher query interface. The top query is:

```
neo4j$ MATCH(b:Blog)←[:wrote]-(a)-[:wrote]→(p:Post{techcrunch_post:null}) RETURN a.name
```

The result is: (no changes, no records). The query completed after 6946 ms.

The bottom query is:

```
1 MATCH (a)-[:wrote]→(:Post{techcrunch_post:null})
2 WHERE (:Blog)←[:wrote]-(a)
3 RETURN a.name
```

The result is: (no changes, no records). The query completed after 7 ms.

### PostgreSQL:

The image shows the PostgreSQL SQL query interface. The query is:

```
8
9 SELECT a
10 FROM authors as a
11 JOIN posts as b on a.name = b.author_name and b.techcrunch_post is false
12 JOIN authors as a1 on a.name = a1.name
13 JOIN posts as p on a1.name = p.author_name and p.techcrunch_post is true
14
```

The interface shows tabs for Data Output, Messages, and Notifications. The Data Output tab is active, showing a table with one row:

a
authors

The status bar at the bottom shows: Total rows: 0 of 0, Query complete 00:00:00.247.

Разлика: -6 699 мс, 240 мс.

6. Автори кои си коментирале на свој пост на истиот ден кога го напишале

#### Neo4j:

```
1 MATCH(a:Author)-[:wrote]→(p:Post)←[c:commented]-(a)
2 WHERE date(p.date)=date(c.date)
3 RETURN distinct a.name
```

	a.name
1	"Michael Arrington"
2	"Devin Coldewey"
3	"MG Siegler"
4	"Andy Brett"
5	"Erick Schonfeld"
6	"Robin Wauters"
7	

Started streaming 57 records after 1 ms and completed after 438 ms.

#### PostgreSQL:

```
13
14 SELECT distinct(p.author_name)
15 FROM posts as p
16 JOIN comments as c on c.author_name = p.author_name
17 WHERE c.post_url = p.url and date(p.date) = date(c.date)
```

	author_name character varying
1	Andy Brett
2	Arden Pennell
3	Asad Akbar
4	Blake Robinson
5	Brian Solis
6	Chris Velazco
7	Dan Kimerling
8	Dan Romero
9	David Diaz
10	Devin Coldewey

Total rows: 57 of 57    Query complete 00:00:01.043    ✓ PostgreSQL 15/Bloggers - Database connected. ✕

Разлика: 605 мс

## 7. Постови кои се линкуваат меѓусебно

### Neo4j:

```
1 MATCH(p:Post)-[:links *2]→(p)
2 RETURN p.title
```

	p.title
1	"Six New Years Resolutions For Apple And The iPhone In 2010"
2	"The Apple Tablet: Will It Be Called iSlate, iGuide, Or Something Else?"

Started streaming 2 records after 1 ms and completed after 260 ms.

### PostgreSQL:

```
8 SELECT p.url
9 FROM posts as p
10 JOIN inlinks as il ON il.post_url=p.url
11 JOIN posts as p1 on il.linked_post_url = p1.url
12 JOIN inlinks as il1 ON il1.post_url=p1.url AND il1.linked_post_url = p.url
13
```

Data Output   Messages   Notifications

	url [PK] character varying
1	<a href="http://www.mobilecrunch.com/2010/01/01/five-new-years-resolutions-for-apple-and-the-iphone-in-20...">http://www.mobilecrunch.com/2010/01/01/five-new-years-resolutions-for-apple-and-the-iphone-in-20...</a>
2	<a href="http://techcrunch.com/2009/12/29/apple-tablet-islate-iguide/">http://techcrunch.com/2009/12/29/apple-tablet-islate-iguide/</a>

Total rows: 2 of 2   Query complete 00:00:00.288

Разлика 28 мс

## 8. Постови кои во содржината го содржат името на авторот

### Neo4j:

```
1 MATCH(p:Post)←[:wrote]-(a:Author)
2 WHERE p.title CONTAINS a.name
3 RETURN p.title, a.name
```

Table

p.title

"will skye minutes become fungible?"

15 "stumbling through //engtech at random (or any wordpress.com blog)"

16 "del.zio.si x lafra [mag1]: aol copia yahoo, il mestiere di"

17 "fitnessmantra weekend: see how much you are eating; even before"

18 "a yahoo legyilkolja saj"

19 "Van apareciendo nuevos portales relacionados con lo que nos ocupa II"

20 "chubby at foo camp: one year later"

Started streaming 130 records after 1 ms and completed after 310 ms.

### PostgreSQL:

```
10 SELECT p.title, p.author_name
11 FROM posts as p
12 WHERE p.content LIKE CONCAT('%',p.author_name,'%')
13
14
```

Data Output Messages Notifications

	title character varying	author_name character varying
1	GeeknRolla To Launch Startups, London, April 20	Mike Butcher
2	Hello, iPad. Hello, Cloud 2.	Marc Benioff
3	Microsoft&#8217;s Scott Guthrie on Silverlight and Windows Phone	Steve Gillmor
4	A Fix for Discrimination: Follow the Indian Trails	Vivek Wadhwa
5	What's Better: Saving the World or Building Another Facebook app?	Vivek Wadhwa
6	Silicon Valley: You and Some of Your VC&#039;s have a Gender Problem	

Total rows: 218 of 218 Query complete 00:00:00.204

✓ Successfully run. Total query runtime: 204 msec. 21

Разлика: -106 мс



## Neo4j:

## PostgreSQL:

Разлика: 2 593 мс

## 10. Блогери и вкупни коментари на сите нивни постови во опаѓачки редослед

Neo4j:

```
1 MATCH (a:Blogger)-[:wrote]-(p:Post)←[:commented]-(c)
2 WITH a, count(c) as comment_count
3 ORDER BY comment_count DESC
4 RETURN a.name, comment_count
```

	a.name	comment_count
1	"Michael Arrington"	245674
2	"Erick Schonfeld"	95017
3	"MG Siegler"	71739
4	"Jason Kincaid"	61966
5	"Robin Wauters"	45271
6	"Duncan Riley"	32708
7		

Started streaming 102 records in less than 1 ms and completed after 90 ms.

PostgreSQL:

```
9
10 SELECT name, comment_count
11 FROM (
12     SELECT a.name as name, count(c) as comment_count
13     FROM authors as a
14     JOIN posts as p on p.author_name = a.name and a.techcrunch_blogger is true
15     JOIN comments as c on c.post_url= p.url
16     GROUP BY a.name
17 ) as bloggers
18 ORDER BY comment_count DESC
19
```

	name [PK] character varying	comment_count bigint
1	Michael Arrington	245674
2	Erick Schonfeld	95017
3	MG Siegler	71739
4	Jason Kincaid	61966
5	Robin Wauters	45271
6	Duncan Riley	32708

Total rows: 102 of 102    Query complete 00:00:01.602    ✓ Successfully run. Total query runtime: 1 secs 602 ms

Разлика: 1 512 мс

## **Заклучок:**

За поголемиот дел од прашалниците важи дека Neo4j побрзо дава резултат, иако има и исклучоци.

Neo4j е впечатливо побрза во случајот 9 и 10, или случаите каде што во Neo4j се избегнува употреба на вгнездено квери.

Од 2. и 3. Може да се забележи дека додавањето на дополнителни лабели на ентитетите во Neo4j може значително да ги забрза кверињата, додека во релациона база нема таква можност.

Во 4 Neo4j повторно има подобри резултати и многу поедноставен начин на пишување на квери, отколку во PostgreSQL каде треба да се направат повеќе спојувања на табели.

Во 1 и 8 се работи за квериња кои се концентрираат на една табела, и нема голема разлика во перформансите во двете бази.

Во повеќето квериња каде има поврзување помеѓу различни табели Neo4j има подобри резултати, иако нема премногу голема разлика.

Од 2 и 5 може да се забележи дека кверињата во кои се работи за различни ентитети со иста лабела кои се во ист тип на релација со некој ентитет со различна лабела имаат многу полоши перформанси. Но, промена на формата на кварино може да ја промени брзината на извршување многу пати.

Кога пак се работи за прашалници каде се кверираат ентитети од ист тип поврзани со истата релација, Neo4j има подобри перформанси(7 и 4).

Како и да е, прашалниот јазик на Neo4, Cypher има многу полесна синтакса за користење.

Од овој експеримент може да се заклучи дека Neo4j има подобри перформанси кога пребарувањето во базата се врши на база на релации, но потребно е да се внимава каков тип на прашалници се најчести.